

In [25]:

```
# Load CIFAR-10 data set
from keras.datasets import cifar10
(X_train, y_train), (X_test, y_test) = cifar10.load_data()
```

In [26]:

```
# Show examples from each class
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd

num_classes = len(np.unique(y_train))
class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
fig = plt.figure(figsize=(8,3))
for i in range(num_classes):
    ax = fig.add_subplot(2, 5, 1 + i, xticks=[], yticks=[])
    ax.set_title(class_names[i])
    idx = np.where(y_train[:] == i)[0]
    features_idx = X_train[idx, :]
    rnd_img = np.random.randint(features_idx.shape[0])
    im = np.transpose(features_idx[rnd_img, :], (0, 1, 2))
    plt.imshow(im)
plt.show()
```



In [27]:

```
# Data pre-processing
X_train = X_train.astype('float32')
X_test = X_test.astype('float32')
X_train /= 255.0
X_test /= 255.0

from keras.utils import np_utils
Y_train = np_utils.to_categorical(y_train, num_classes)
Y_test = np_utils.to_categorical(y_test, num_classes)
```

In [28]:

```
def plot_model_history(model_history):
    fig, axs = plt.subplots(1,2,figsize=(15,5))
    # Summarize history for accuracy
    axs[0].plot(range(1,len(model_history.history['accuracy'])+1),model_history.history['accuracy'])
    axs[0].plot(range(1,len(model_history.history['val_accuracy'])+1),model_history.history['val_accuracy'])
    axs[0].set_ylim(0, 1)
    axs[0].set_title('Model Accuracy')
    axs[0].set_ylabel('Accuracy')
    axs[0].set_xlabel('Epoch')
    axs[0].set_xticks(np.arange(1,len(model_history.history['accuracy'])+1,step=len(model_history.history['accuracy'])))
```

```

axs[0].legend(['train', 'val'], loc='best')
# summarize history for loss
axs[1].plot(range(1,len(model_history.history['loss'])+1),model_history.history[
axs[1].plot(range(1,len(model_history.history['val_loss'])+1),model_history.hist
axs[1].set_title('Model Loss')
axs[1].set_ylabel('Loss')
axs[1].set_xlabel('Epoch')
axs[1].set_xticks(np.arange(1,len(model_history.history['loss'])+1,step=len(mode
axs[1].legend(['train', 'val'], loc='best')
plt.show()

```

In [29]:

```

# Feed forward dense model
# Here it is not allowed to use convolutional layers
# You may use any regularizacoin (see class slides)

from keras.models import Sequential
from keras.layers import Dense, Dropout, Activation, Flatten, BatchNormalization
from keras import optimizers
from keras.callbacks import EarlyStopping
from keras.regularizers import l2

learning_rate=0.001
epochs=1000
batch_size=512
es = EarlyStopping(monitor='val_loss', mode='min', verbose=1, patience=int(epochs*0.
p_dropout=0.2

# model.add(Dense(500, activation="relu", kernel_initializer="he_normal"))

model = Sequential()
model.add(Flatten(input_shape=(32, 32, 3)))
model.add(Dense(200, activation="relu"))
model.add(Dropout(rate=p_dropout))
model.add(BatchNormalization())
model.add(Dense(200, activation="relu"))
model.add(Dropout(rate=p_dropout))
model.add(BatchNormalization())
model.add(Dense(150, activation="relu"))
model.add(Dropout(rate=p_dropout))
model.add(BatchNormalization())
model.add(Dense(150, activation="relu"))
model.add(Dropout(rate=p_dropout))
model.add(BatchNormalization())
model.add(Dense(100, activation="relu"))
model.add(Dropout(rate=p_dropout))
model.add(BatchNormalization())
model.add(Dense(num_classes, activation='softmax'))

opt = optimizers.SGD(lr=learning_rate, momentum=0.9, nesterov=True)
# opt = optimizers.Adam(lr=learning_rate, beta_1=0.9, beta_2=0.999)

model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])
model.summary()

```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
flatten_3 (Flatten)	(None, 3072)	0

dense_18 (Dense)	(None, 200)	614600
dropout_15 (Dropout)	(None, 200)	0
batch_normalization_15 (Batch Normalization)	(None, 200)	800
dense_19 (Dense)	(None, 200)	40200
dropout_16 (Dropout)	(None, 200)	0
batch_normalization_16 (Batch Normalization)	(None, 200)	800
dense_20 (Dense)	(None, 150)	30150
dropout_17 (Dropout)	(None, 150)	0
batch_normalization_17 (Batch Normalization)	(None, 150)	600
dense_21 (Dense)	(None, 150)	22650
dropout_18 (Dropout)	(None, 150)	0
batch_normalization_18 (Batch Normalization)	(None, 150)	600
dense_22 (Dense)	(None, 100)	15100
dropout_19 (Dropout)	(None, 100)	0
batch_normalization_19 (Batch Normalization)	(None, 100)	400
dense_23 (Dense)	(None, 10)	1010
=====		
Total params: 726,910		
Trainable params: 725,310		
Non-trainable params: 1,600		

In [30]:

```
# Training
import time
start = time.time()
history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=epochs, verbose=0)
end = time.time()
```

```
Epoch 1/1000
98/98 [=====] - 2s 13ms/step - loss: 2.8288 - accuracy: 0.1219 - val_loss: 2.2182 - val_accuracy: 0.2020
Epoch 2/1000
98/98 [=====] - 1s 9ms/step - loss: 2.4514 - accuracy: 0.1824 - val_loss: 2.0735 - val_accuracy: 0.2550
Epoch 3/1000
98/98 [=====] - 1s 9ms/step - loss: 2.2851 - accuracy: 0.2159 - val_loss: 1.9689 - val_accuracy: 0.2917
Epoch 4/1000
98/98 [=====] - 1s 10ms/step - loss: 2.1777 - accuracy: 0.2423 - val_loss: 1.8853 - val_accuracy: 0.3249
Epoch 5/1000
98/98 [=====] - 1s 11ms/step - loss: 2.0893 - accuracy: 0.2611 - val_loss: 1.8571 - val_accuracy: 0.3289
Epoch 6/1000
98/98 [=====] - 1s 11ms/step - loss: 2.0402 - accuracy: 0.2742 - val_loss: 1.8236 - val_accuracy: 0.3469
Epoch 7/1000
98/98 [=====] - 1s 10ms/step - loss: 2.0037 - accuracy: 0.28
```

```

548 - val_loss: 1.5638 - val_accuracy: 0.4862
Epoch 333/1000
98/98 [=====] - 1s 11ms/step - loss: 0.9789 - accuracy: 0.6
465 - val_loss: 1.5065 - val_accuracy: 0.5032
Epoch 334/1000
98/98 [=====] - 1s 10ms/step - loss: 0.9628 - accuracy: 0.6
582 - val_loss: 1.5324 - val_accuracy: 0.4995
Epoch 335/1000
98/98 [=====] - 1s 10ms/step - loss: 0.9667 - accuracy: 0.6
546 - val_loss: 1.5191 - val_accuracy: 0.5047
Epoch 00335: early stopping

```

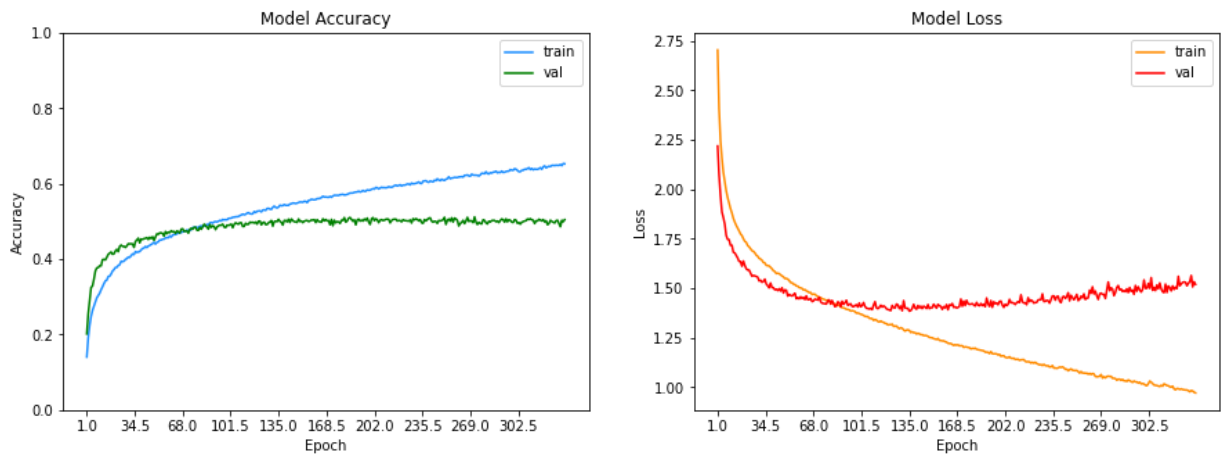
Definición de una red convolucional multicapa

In [31]:

```

plot_model_history(history)
print("Training MLP took " + str(end - start) + " seconds")

```



Training MLP took 335.1605944633484 seconds

In [33]:

```

loss, acc = model.evaluate(X_test, Y_test, verbose=0)
print('Test loss:', loss)
print('Test accuracy:', acc)

```

```

Test loss: 1.503757273864746
Test accuracy: 0.5597000050544739

```