

# The German Traffic Sign Benchmark

Student Name 1: Alberto Miño

Student Name 2: Adrian Michelena

```
In [ ]: # Download the data base
# !wget -c http://www.dia.fi.upm.es/~lbaumela/FULLIJCNN2013.zip
# !unzip FullIJCNN2013.zip
from google.colab import drive
drive.mount('/gdrive', force_remount=True)

# drivePrefix = "/gdrive/My Drive/Colab Notebooks/MUIA-ComputerVision/P4/dataset/"
!unzip -qq -u "/gdrive/My Drive/Colab Notebooks/MUIA-ComputerVision/P4/dataset/FULLI
```

Mounted at /gdrive

```
In [ ]: def plot_model_history(model_history):
    fig, axs = plt.subplots(1,2,figsize=(15,5))
    # Summarize history for accuracy
    axs[0].plot(range(1,len(model_history.history['accuracy'])+1),model_history.history['accuracy'])
    axs[0].plot(range(1,len(model_history.history['val_accuracy'])+1),model_history.history['val_accuracy'])
    axs[0].set_ylim(0, 1)
    axs[0].set_title('Model Accuracy')
    axs[0].set_ylabel('Accuracy')
    axs[0].set_xlabel('Epoch')
    axs[0].set_xticks(np.arange(1,len(model_history.history['accuracy'])+1,step=len(model_history.history['accuracy'])))
    axs[0].legend(['train', 'val'], loc='best')
    # summarize history for loss
    axs[1].plot(range(1,len(model_history.history['loss'])+1),model_history.history['loss'])
    axs[1].plot(range(1,len(model_history.history['val_loss'])+1),model_history.history['val_loss'])
    axs[1].set_title('Model Loss')
    axs[1].set_ylabel('Loss')
    axs[1].set_xlabel('Epoch')
    axs[1].set_xticks(np.arange(1,len(model_history.history['loss'])+1,step=len(model_history.history['loss'])))
    axs[1].legend(['train', 'val'], loc='best')
    plt.show()
```

```
In [ ]: import numpy as np
import cv2
import pandas as pd

IMG_HEIGHT = 600
SIGN_SIZE = (224, 224)

# Function for reading the images
def readImages(rootpath, images_range, signs_range):
    '''Reads traffic sign data for German Traffic Sign Recognition Benchmark.
    Arguments: path to the traffic sign data, for example 'FullIJCNN2013'
    Returns: list of images, list of corresponding labels'''
    images = {} # original image
    scales = {} # original scale
    for num in images_range:
        filename = rootpath + '/' + "{:05d}".format(num) + '.ppm'
        img = cv2.imread(filename, cv2.IMREAD_COLOR)
        scale = IMG_HEIGHT / float(img.shape[0])
```

```

img_resized = cv2.resize(img, (int(img.shape[1]*scale),int(img.shape[0]*scale)))
images.setdefault(filename,[]).append(img_resized)
scales.setdefault(filename,[]).append(scale)

files = [] # filenames
signs = [] # traffic sign image
bboxes = [] # corresponding box detection
labels = [] # traffic sign type
data = np.genfromtxt(rootpath + '/' + 'gt.txt', delimiter=';', dtype=str, usecols=(0,1,2,3,4,5))
for elem in signs_range:
    filename = rootpath + '/' + data[elem][0]
    img = images.get(filename)[0]
    scale = scales.get(filename)[0]
    bbox = np.array([int(data[elem][1]), int(data[elem][2]), int(data[elem][3]), int(data[elem][4])])
    sign = img[int(bbox[1]):int(bbox[3]), int(bbox[0]):int(bbox[2])]
    sign_resized = cv2.resize(sign, SIGN_SIZE)
    files.append(filename)
    signs.append(sign_resized)
    bboxes.append(bbox)
    labels.append(data[elem][5])
return images, files, signs, bboxes, labels

```

```

In [ ]: # The German Traffic Sign Recognition Benchmark
train_images, train_files, train_signs, train_bboxes, train_labels = readImages('FullIJCN')
test_images, test_files, test_signs, test_bboxes, test_labels = readImages('FullIJCN')

```

```

In [ ]: import matplotlib.pyplot as plt
%matplotlib inline

# Show examples from each class
class_names = np.unique(train_labels)
num_classes = len(class_names)
fig = plt.figure(figsize=(8,8))
for i in range(num_classes):
    ax = fig.add_subplot(6, 9, 1 + i, xticks=[], yticks=[])
    ax.set_title(class_names[i])
    indices = np.where(np.isin(train_labels, class_names[i]))[0]
    plt.imshow(cv2.cvtColor(train_signs[int(np.random.choice(indices, 1))], cv2.COLOR_BGR2RGB))
plt.show()

```

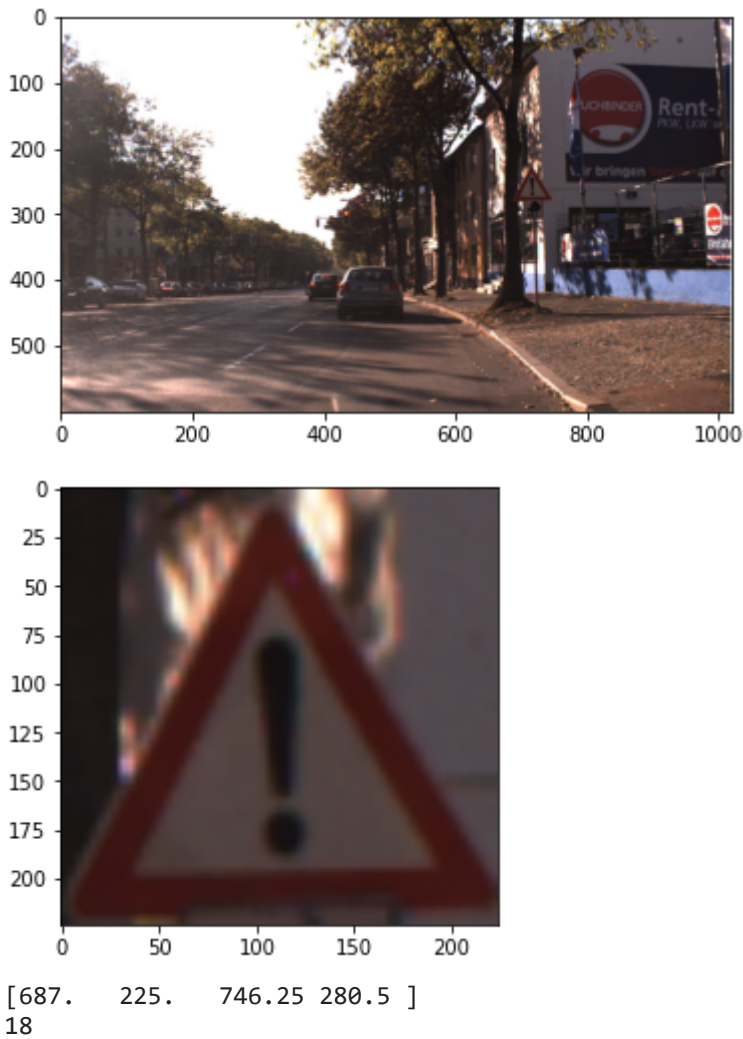


```
In [ ]: from sklearn.utils import shuffle
train_files, train_signs, train_bboxes, train_labels = shuffle(train_files, train_si
plt.imshow(cv2.cvtColor(train_images.get(train_files[0])[0], cv2.COLOR_BGR2RGB))
plt.show()
plt.imshow(cv2.cvtColor(train_signs[0], cv2.COLOR_BGR2RGB))
plt.show()
print(train_bboxes[0])
print(train_labels[0])

# Data pre-processing
tr_signs = np.array(train_signs)[0:600]
tr_labels = np.array(train_labels)[0:600]
va_signs = np.array(train_signs)[600:852]
va_labels = np.array(train_labels)[600:852]
te_signs = np.array(test_signs)
te_labels = np.array(test_labels)

tr_signs = tr_signs.astype('float32')
va_signs = va_signs.astype('float32')
te_signs = te_signs.astype('float32')
tr_signs /= 255.0
va_signs /= 255.0
te_signs /= 255.0

from keras.utils import np_utils
tr_labels = np_utils.to_categorical(tr_labels, num_classes)
va_labels = np_utils.to_categorical(va_labels, num_classes)
te_labels = np_utils.to_categorical(te_labels, num_classes)
```



```
In [ ]: # Tensorboard
from time import time
from keras.callbacks import TensorBoard
tensorboard = TensorBoard(log_dir='logs/{}'.format(time()))
```

```
In [ ]: # Data augmentation
from keras.preprocessing.image import ImageDataGenerator
datagen = ImageDataGenerator(
    # featurewise_center=True,
    # featurewise_std_normalization=True,
    rotation_range=10.,
    width_shift_range=0.1,
    height_shift_range=0.1,
    horizontal_flip=False,
    # zoom_range=0.15,
    shear_range=0.1,
    fill_mode='nearest')

datagen.fit(tr_signs)
```

## Assignment 1: Multi-Layer Perceptron

```
In [ ]: # Convolutional Neural Network (CNN)
# Here you are allowed to use convolutional layers
# You may use also any regularization (see class slides)
```

```

from keras.models import Sequential
from keras.layers import Dense, Activation, Flatten, Dropout, BatchNormalization
from keras import optimizers
from keras.callbacks import EarlyStopping
from keras.regularizers import l2
from keras.layers.convolutional import Conv2D, MaxPooling2D
import keras.backend as K
from keras.callbacks import ModelCheckpoint

checkpoint_path = "/gdrive/My Drive/Colab Notebooks/MUIA-ComputerVision/P4/Alberto_T
checkpoint_callback = ModelCheckpoint(
    checkpoint_path, monitor='val_accuracy', verbose=1, save_weights_only=True,
    # Save weights, every epoch.
    save_freq='epoch', mode='auto', save_best_only=True)

learning_rate=0.001
epochs=1000
batch_size=32
es = EarlyStopping(monitor='val_loss', mode='auto', verbose=1, patience=int(epochs*0
p_dropou_layert=[0,0,0.2,0.3,0.5,0.5,0.7]
i=0
d_augm=1

model = Sequential()
model.add(Conv2D(filters=16, kernel_size=(3, 3), activation='relu', padding='same', i
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(p_dropou_layert[i]))
i+=1
model.add(Conv2D(filters=32, kernel_size=(3, 3), activation='relu', padding='same'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(p_dropou_layert[i]))
i+=1
model.add(Conv2D(filters=64, kernel_size=(3, 3), activation='relu', padding='same'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(p_dropou_layert[i]))
i+=1
model.add(Conv2D(filters=128, kernel_size=(3, 3), activation='relu', padding='same'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(p_dropou_layert[i]))
i+=1
model.add(Conv2D(filters=256, kernel_size=(3, 3), activation='relu', padding='same'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(BatchNormalization())
model.add(Dropout(p_dropou_layert[i]))
i+=1
model.add(Conv2D(filters=512, kernel_size=(3, 3), activation='relu', padding='same'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(BatchNormalization())
model.add(Dropout(p_dropou_layert[i]))
i+=1
model.add(Conv2D(filters=512, kernel_size=(3, 3), activation='relu', padding='same'))
model.add(MaxPooling2D(pool_size=(1, 1)))
model.add(BatchNormalization())
model.add(Dropout(p_dropou_layert[i]))

model.add(Conv2D(filters=num_classes, kernel_size=(3, 3), padding='valid'))
model.add(Flatten())

```

```

model.add(Activation('softmax'))

# opt = optimizers.SGD(lr=learning_rate, momentum=0.9, nesterov=True)
opt = optimizers.Adam(lr=learning_rate, beta_1=0.9, beta_2=0.999)

model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])
model.summary()

```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
=====		
conv2d_8 (Conv2D)	(None, 224, 224, 16)	448
max_pooling2d_7 (MaxPooling2)	(None, 112, 112, 16)	0
dropout_7 (Dropout)	(None, 112, 112, 16)	0
conv2d_9 (Conv2D)	(None, 112, 112, 32)	4640
max_pooling2d_8 (MaxPooling2)	(None, 56, 56, 32)	0
batch_normalization_6 (Batch Normalization)	(None, 56, 56, 32)	128
dropout_8 (Dropout)	(None, 56, 56, 32)	0
conv2d_10 (Conv2D)	(None, 56, 56, 64)	18496
max_pooling2d_9 (MaxPooling2)	(None, 28, 28, 64)	0
batch_normalization_7 (Batch Normalization)	(None, 28, 28, 64)	256
dropout_9 (Dropout)	(None, 28, 28, 64)	0
conv2d_11 (Conv2D)	(None, 28, 28, 128)	73856
max_pooling2d_10 (MaxPooling2)	(None, 14, 14, 128)	0
batch_normalization_8 (Batch Normalization)	(None, 14, 14, 128)	512
dropout_10 (Dropout)	(None, 14, 14, 128)	0
conv2d_12 (Conv2D)	(None, 14, 14, 256)	295168
max_pooling2d_11 (MaxPooling2)	(None, 7, 7, 256)	0
batch_normalization_9 (Batch Normalization)	(None, 7, 7, 256)	1024
dropout_11 (Dropout)	(None, 7, 7, 256)	0
conv2d_13 (Conv2D)	(None, 7, 7, 512)	1180160
max_pooling2d_12 (MaxPooling2)	(None, 3, 3, 512)	0
batch_normalization_10 (Batch Normalization)	(None, 3, 3, 512)	2048
dropout_12 (Dropout)	(None, 3, 3, 512)	0
conv2d_14 (Conv2D)	(None, 3, 3, 512)	2359808
max_pooling2d_13 (MaxPooling2)	(None, 3, 3, 512)	0
batch_normalization_11 (Batch Normalization)	(None, 3, 3, 512)	2048

dropout_13 (Dropout)	(None, 3, 3, 512)	0
conv2d_15 (Conv2D)	(None, 1, 1, 43)	198187
flatten_1 (Flatten)	(None, 43)	0
activation_1 (Activation)	(None, 43)	0
=====		
Total params: 4,136,779		
Trainable params: 4,133,771		
Non-trainable params: 3,008		

In [ ]:

```

if d_augm==1:
    # Fit the model with real time data augmentation
    print("Fitting model with data augmentation")
    start = time()
    data = model.fit(datagen.flow(tr_signs,tr_labels,batch_size=batch_size),
                      epochs=epochs, verbose=2, validation_data=(va_signs, va_l
    end = time()

else:
    # Fit the model with plain dataset
    print("Fitting model")
    start = time()
    data = model.fit(tr_signs, tr_labels, batch_size=batch_size, epochs=epochs, verbos
    end = time()

```

Fitting model with data augmentation

Epoch 1/1000

19/19 - 8s - loss: 5.4695 - accuracy: 0.0800 - val\_loss: 3.8297 - val\_accuracy: 0.0397

Epoch 00001: val\_accuracy improved from -inf to 0.03968, saving model to /gdrive/My Drive/Colab Notebooks/MUIA-ComputerVision/P4/Alberto\_TrafficSignRecognitionCNN2/best\_epoch\_val\_acc.ckpt

Epoch 2/1000

19/19 - 6s - loss: 4.6520 - accuracy: 0.1517 - val\_loss: 4.6501 - val\_accuracy: 0.0000e+00

Epoch 00002: val\_accuracy did not improve from 0.03968

Epoch 3/1000

19/19 - 7s - loss: 4.2820 - accuracy: 0.2067 - val\_loss: 4.6173 - val\_accuracy: 0.0238

Epoch 00003: val\_accuracy did not improve from 0.03968

Epoch 4/1000

19/19 - 6s - loss: 3.7169 - accuracy: 0.2417 - val\_loss: 6.0639 - val\_accuracy: 0.0198

Epoch 00004: val\_accuracy did not improve from 0.03968

Epoch 5/1000

19/19 - 6s - loss: 3.5540 - accuracy: 0.2717 - val\_loss: 5.3848 - val\_accuracy: 0.0198

Epoch 00005: val\_accuracy did not improve from 0.03968

Epoch 6/1000

19/19 - 6s - loss: 3.1447 - accuracy: 0.2933 - val\_loss: 6.1567 - val\_accuracy: 0.0198

Epoch 00006: val\_accuracy did not improve from 0.03968

Epoch 7/1000

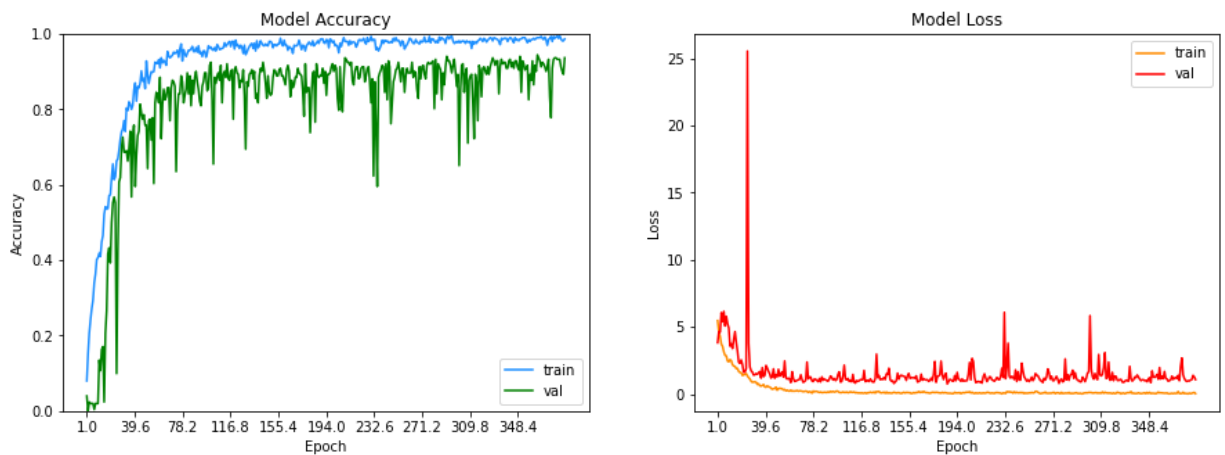
Epoch 386/1000

19/19 - 6s - loss: 0.0400 - accuracy: 0.9867 - val\_loss: 1.0839 - val\_accuracy: 0.9365

Epoch 00386: val\_accuracy did not improve from 0.94444

Epoch 00386: early stopping

```
In [ ]: print("Training CNN took " + str(end - start) + " seconds")
        plot_model_history(data)
```



```
In [ ]: start = time()
        loss, acc = model.evaluate(te_signs, te_labels, verbose=0)
        end = time()
        print('CNN took ' + str(end - start) + ' seconds')
        print('For final weights configuration:\n\tTest loss: ' + str(loss) + ' - Accuracy: ' + str(acc))
```

CNN took 0.2529294490814209 seconds

For final weights configuration:

Test loss: 0.5962536334991455 - Accuracy: 0.9639889001846313

```
In [ ]: model.load_weights(checkpoint_path)
        start = time()
        loss, acc = model.evaluate(te_signs, te_labels, verbose=0)
        end = time()
        print('CNN took ' + str(end - start) + ' seconds')
        print('For best validation accuracy weights configuration found in training:\n\tTest loss: ' + str(loss) + ' - Accuracy: ' + str(acc))
```

CNN took 0.24784636497497559 seconds

For best validation accuracy weights configuration found in training:

Test loss: 0.638451337814331 - Accuracy: 0.9667590260505676